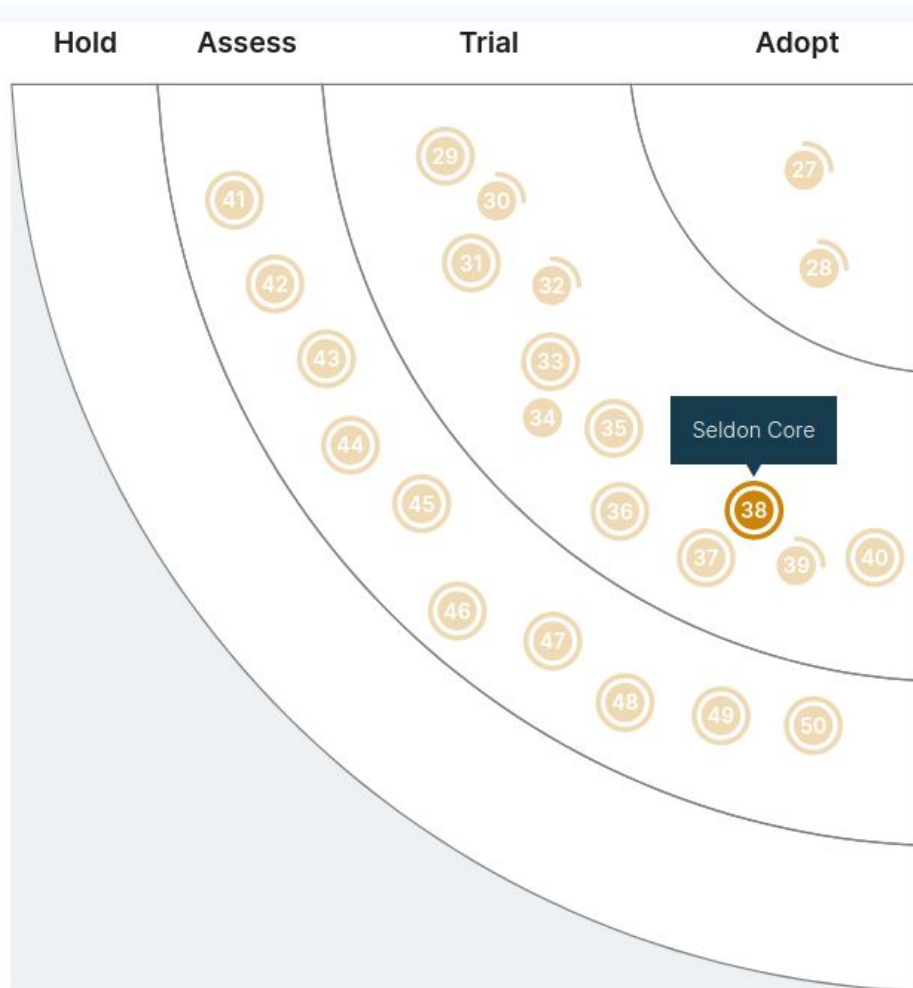


Seldon Core

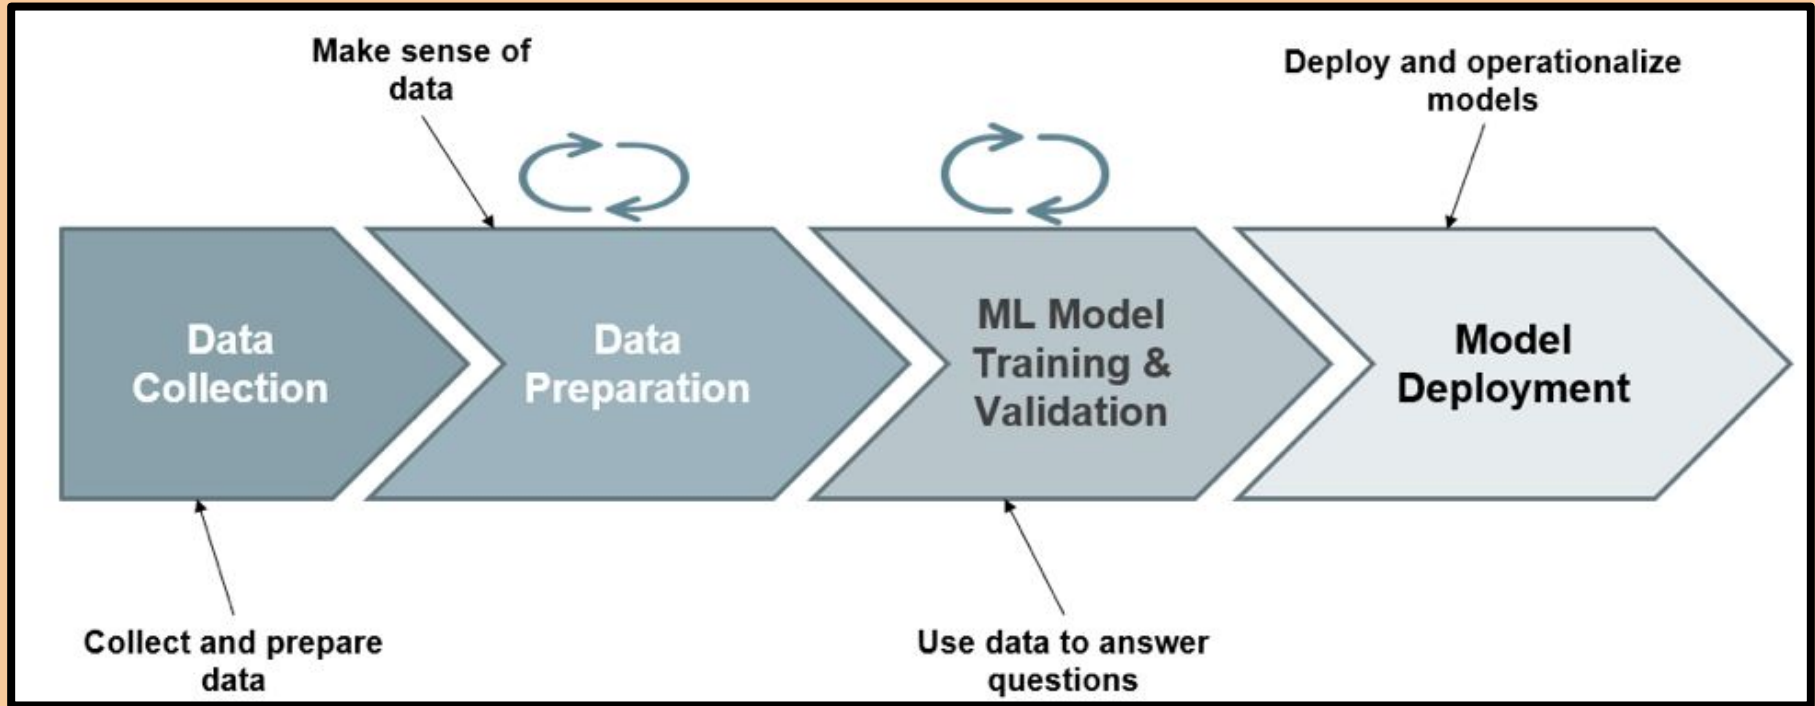
38. Seldon Core

Seldon Core is an open-source platform to package, deploy, monitor and manage machine learning models in [Kubernetes](#) clusters. With out-of-the-box support for several machine-learning frameworks, you can easily containerize your models using [prepackaged inference servers](#), [custom inference servers](#) or [language wrappers](#). With distributed tracing through [Jaeger](#) and model explainability via [Alibi](#), Seldon Core addresses several last-mile delivery challenges with machine learning deployments, and our data teams like it.

View blip history >



Machine Learning overview



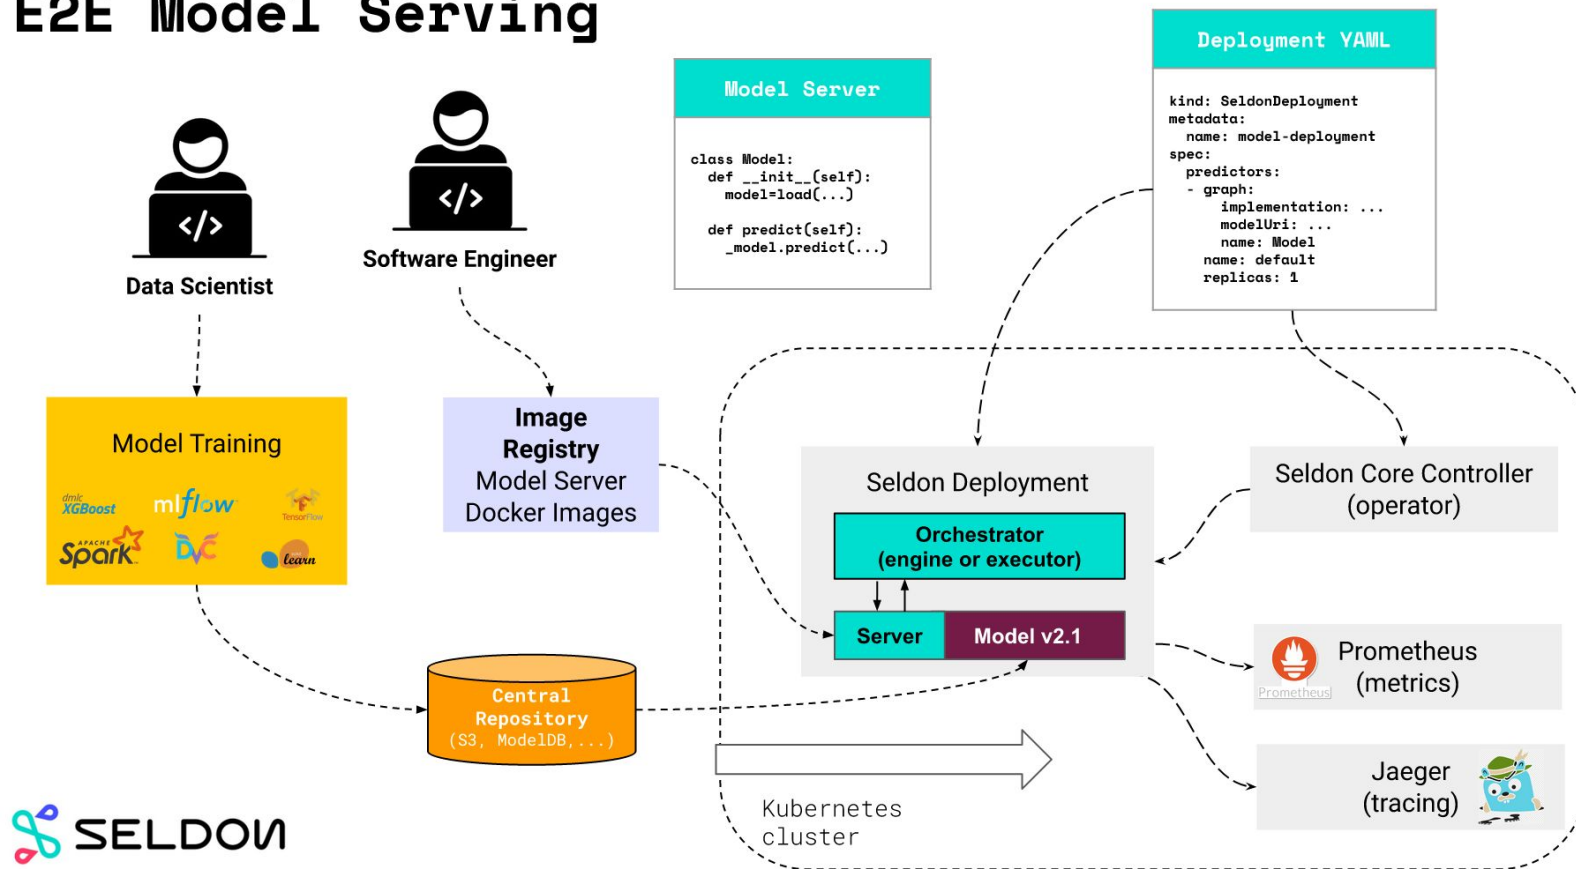
Kubernetes overview

- Gives you a way to centrally manage deploying containers to multiple nodes.
- You can tell the Kubernetes API that you want **3 objects**:
 - A **Deployment** of a Node.js container with 2 replicas
 - A **Service** to map port 3000 on the containers to port 80 internally to other apps inside Kubernetes at <http://pottery-barn-app>
 - An **Ingress** to tell a load balancer that you'd like to expose your Service to the world at <https://bobert-family-pottery-barn.com>
- It's declarative — you tell it the definition of what you want, and the controllers inside Kubernetes make it happen.

Seldon Core overview

- Gives you a way to deploy your trained ML model, with extra features.
- Automatically generates:
 - OpenAPI docs (Swagger)
 - Prometheus metrics
 - Jaeger traces
- Provides premade servers for popular frameworks like scikit-learn, Tensorflow, and XGBoost.
 - And you then just provide the location of the model in blob storage (like s3).
- Includes the ability to set up pre-processors for the data before the model server receives it.

E2E Model Serving



Questions?